


Specializing a released model vs. scaling it up

■ Kev (released) ■ RadKev (specialized)

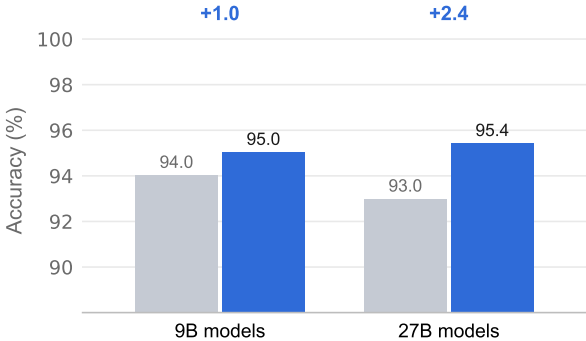
a All human-labeled questions

$\Delta = \text{RadKev} - \text{Kev}$



b Radiology questions only

$\Delta = \text{RadKev} - \text{Kev}$

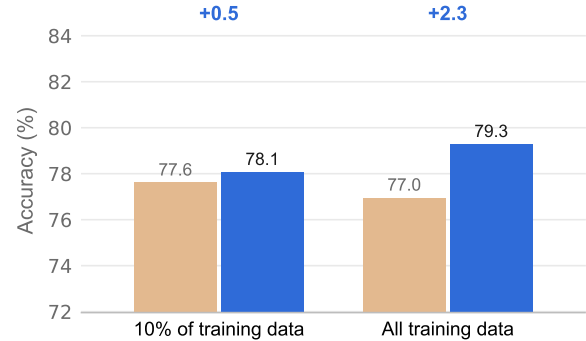


Fine-tuning a 9B model from Kev-9B vs. from its base model

■ started from Qwen3.5-9B-Base ■ started from Kev-9B

c Accuracy after fine-tuning

All human-labeled questions; $\Delta = \text{Kev-9B start} - \text{base-model start}$



d General ability outside radiology

Kev's transfer suite, relative to the released Kev-9B

